

MIT EECS 6.815/6.865: Assignment 5:
Resampling, Warping and Morphing

Due Wednesday October 20 at 9pm

1 Summary

This problem set has several questions for extra credit. Feel free to attempt them, but do the main problems first. The maximum extra credit is 10%. The last section of this is submitting a photo of yourself, which might take some time depending on what type of person you are. So don't put it off!

- Scaling using nearest-neighbor and bilinear reconstruction
- Scaling using bicubic and Lanczos methods (6.865 only)
- Image warping according to one pair of segments
- Image warping according to two lists of segments, using weighted warping
- Image morphing
- Take a photo of yourself

2 Resampling

In this section, we will rescale images, starting with simple transformations and naïve reconstructions. See Figure 1 for examples of scaling with various methods, together with a small crop of the resulting image to highlight the differences.

2.1 Basic scaling with nearest-neighbor


The first operation we consider is the re-scaling of an image by a global scale factor s . If $s > 1$, the operation will enlarge the input. If $0 < s < 1$, it will shrink the input.

To implement this operation, we need to create a new **Image** object that is s times larger (resp. $\frac{1}{s} \times$ smaller if $s < 1$) than the original in both width and height. Use `floor()` to get an integral size for the new image.


We now have to fill the pixel values of this new image from those of the input: this is called re-sampling. For each pixel in the new image, we look up the color value in the input image at a location that corresponds to the inverse transform (in this case a scaling by $\frac{1}{s}$). In general this location will not be on grid and we'll have to estimate the color at this location using information from the neighboring pixels.




(a) Original Image



(b) Scaling by $3.5\times$ with nearest neighbor interpolation using `scaleNN`



(c) Scaling by $3.5\times$ with bilinear interpolation using `scaleLin`



(d) Scaling by $3.5\times$ with bicubic interpolation

Figure 1: Basic Image scaling.

The simplest technique to sample the new value is called **nearest-neighbor re-sampling**: we round-off the real coordinates to the nearest integers and use the input's color at this new location to be the value of the current output pixel.

1 Implement the `Image scaleNN(const Image &im, float scale)` function in `basicImageManipulation.cpp`. This function should create a new image that is `scale` times the size of the input using the nearest-neighbor re-sampling method.

2.2 Scaling with bilinear interpolation

Nearest-neighbor re-sampling creates blocky artifacts and pixelated results. We will address this using a better reconstruction based on bilinear interpolation. For this, we consider the four pixels immediately around the computed real coordinates and perform two linear interpolations. We first linearly interpolate along x the colors of the top and bottom pairs of pixels. Then we interpolate these two values along y to get the final sample. The interpolation weight are driven by the distance from the corners (see Fig. 2).




Figure 2: Bilinear interpolation. We sample the value at location X (non integer coordinates) by interpolating the neighboring pixels' values. First, linearly interpolate the values of A and B along x . Do the same for D and C . This gives you pixel values at E and F . Linearly interpolate these two values along y to obtain X 's colors.

2.a Implement `float interpolateLin(const Image &im, float x, float y, int z, bool clamp)` in `basicImageManipulation.cpp`. This function takes floating point coordinates and performs bilinear reconstruction on the given channel z . Don't forget to use smart accessors to make sure you can handle coordinates outside the image.

2.b Next, write an image scaling function `Image scaleLin(const Image &im, float scale)` in `basicImageManipulation.cpp` that rescale using linear interpolation by calling `interpolateLin` where appropriate. Use `clamp = true`.

2.3 Bicubic and Lanczos (required for 6.865, 5% Extra credit for 6.815)

You can obtain a better interpolation by considering a larger pixel footprint and using smarter weights, such as that given by a bicubic or Lanczos functions.

You may think of resampling from the view point of convolution to better understand how bicubic/Lanczos resampling kernels work. For example, bilinear interpolation for $3\times$ upscaling would be equivalent to first scaling the image naïvely (by only taking the existing pixels from the source image and keeping the rest as blank) and then applying a 5×5 convolution kernel (how does the kernel look like? A square pyramid) to the naïvely scaled image. **Hint:** This means that an easy implementation of bicubic and Lanczos resampling is to follow the steps of convolution and use *analytical* weights $k(x)$ instead of numerically discretized ones. And bilinear/bicubic/Lanczos resampling is separable, which means you can multiply horizontal and vertical weights $k(x)$ and $k(y)$ to get the 2D weight.

3 Implement the `Image scaleCubic(const Image &im, float scale, float B, float C)` function in `basicImageManipulation.cpp`. This function should create a new image that is `scale` times the size of the input using a bicubic filter kernel. We will use the kernel parametrization from Mitchell and Netravali (“Reconstruction Filters in Computer Graphics”, Mitchell and Netravali 1988):

$$k(x) = \frac{1}{6} \begin{cases} (12 - 9B - 6C)|x|^3 + (-18 + 12B + 6C)|x|^2 + (6 - 2B) & \text{if } |x| < 1 \\ (-B - 6C)|x|^3 + (6B + 30C)|x|^2 + (-12B - 48C)|x| + (8B + 24C) & \text{if } 1 \leq |x| < 2 \\ 0 & \text{otherwise} \end{cases} \quad (1)$$

A nice parameter to test your filter kernel is $B = C = \frac{1}{3}$.

4 Implement the `Image scaleLanczos(const Image &im, float scale, float a)` function in `basicImageManipulation.cpp`. This function should create a new image that is `scale` times the size of the input

using a Lanczos filter kernel. The kernel is:

$$k(x) = \begin{cases} \text{sinc}(x)\text{sinc}(x/a) & \text{if } |x| < a \\ 0 & \text{otherwise} \end{cases} \quad (2)$$

where $\text{sinc}(x) = \sin(\pi x)/(\pi x)$. A nice parameter to test your filter kernel is $a = 3$.

2.4 Rotations (5% extra credit)

5 Implement the function `rotate(const Image &im, float theta)` in `basicImageManipulation.cpp` that rotates an image around its center by θ radians. **Hint:** use your bilinear interpolation function and the center position already present in the starter code.



Figure 3: Rotation with an angle $\theta > 0$ with respect to the center of the image (to get the pixel value at A' , sample from location A in the input image. Last two images: example rotation by $\frac{\pi}{4}$ using `rotate`.

3 Warping and morphing

In what follows, you will implement image warping and morphing according to Beier and Neely's method, which was used for special effects such as those of Michael Jackson's *Black or White* music video (<https://youtu.be/F2AitTPI5U0>).

We highly recommend that you read the provided original article, which is well written and includes important references such as Ghost Busters.

Beier, Thaddeus, and Shawn Neely. "Feature-based image metamorphosis." ACM SIGGRAPH Computer Graphics. Vol. 26. No. 2. ACM, 1992.

The full method for warping and morphing includes a number of technical components and it is critical that you debug them as you implement each individual one. **A copy of the paper is included in the handout.**

3.1 Basic Vector Tools

Warping and morphing geometrically distort an input image. This requires a few vector operation which you'll implement. We provide you a basic `Vec2f` class to represent 2D vectors.

- 6.a In `morphing.cpp`, implement `Vec2f operator+(const Vec2f & a, const Vec2f & b)` to sum two vectors $\mathbf{a} + \mathbf{b}$.
- 6.b In `morphing.cpp`, implement `Vec2f operator-(const Vec2f & a, const Vec2f & b)` that returns the difference $\mathbf{a} - \mathbf{b}$.
- 6.c In `morphing.cpp`, implement `Vec2f operator*(const Vec2f & a, float f)` that implements multiplication by a scalar $f \cdot \mathbf{a}$.
- 6.d In `morphing.cpp`, implement `Vec2f operator/(const Vec2f & a, float f)` that implements division by a scalar \mathbf{a}/f .
- 6.e In `morphing.cpp`, implement `Vec2f dot(const Vec2f &a, const Vec2f &b)` that implements the dot product of two vectors: $a_x * b_x + a_y * b_y$.
- 6.f In `morphing.cpp`, implement `float length(const Vec2f &a)` that returns the length of a vector (in the L^2 sense): $\|\mathbf{a}\| = \sqrt{a_x^2 + a_y^2}$.
- 6.g In `morphing.cpp`, implement `Vec2f perpendicular(const Vec2f &a)` that returns a vector that is perpendicular to \mathbf{a} . **Hint:** either of the two possible directions is fine.

Don't forget to test your functions. We have provided some basic tests for you in `a5_main`.

3.2 Segments

Now that we have some basic tools, we will implement the `Segment` class, which is critical to Beier and Neely's warping. A `Segment` represents a directed line segment \overrightarrow{PQ} . The class holds a copy of the endpoints P and Q , the length of the segment $\|\overrightarrow{PQ}\|$, and a local orthonormal frame (\vec{e}_1, \vec{e}_2) . We define:

$$\vec{e}_1 = \frac{\overrightarrow{PQ}}{\|\overrightarrow{PQ}\|}. \quad (3)$$

\vec{e}_2 completes the orthonormal frame (e.g. \vec{e}_2 is perpendicular to \vec{e}_1), see Figure 4. Both frame vectors and the length of the segment need to be initialized in the class constructor and *can be used directly in its class functions*.




Figure 4: Directed segment and its local coordinate system. Don't forget that u as defined in the paper is not exactly the coordinate in the local frame.

7 In `morphing.cpp`, complete the constructor `Segment::Segment(Vec2f P_, Vec2f Q_)` that creates a segment from its two endpoints and initializes the data structure properly.

Now that our `Segment` class is usable let's implement methods to convert from the global (x, y) coordinates of a point to the local (u, v) coordinates in the reference frame as in Beier and Neely, *e.g.*, equations (1) and (2) in the paper, and the reverse (equation (3)). **A word of caution:** although v exactly corresponds to the second coordinate in the local frame, u is actually rescaled (so read equation (1) carefully).

- 8.a Implement `Vec2f Segment::XtoUV(Vec2f &X)` to compute the (u, v) coordinates of a 2D point $X = (x, y)$ with respect to a segment as described in the paper.
- 8.b Conversely, implement `Vec2f Segment::UVtoX(Vec2f &uv)` to compute the (x, y) global coordinates of a 2D point from its local (u, v) coordinates.
- 8.c Implement the point to segment distance function `float Segment::distance(Vec2f X)` as described in section 3.3 of the paper.

Test these methods thoroughly as they will be used in the warping and morphing functions in the later part of the problem set. The function `testSegment` in `a4.main.cpp` is great place to write simple test cases.

3.3 Warping according to one pair of segments

Now that we have a functional `Segment` class and an interface to specify line segments, let the fun begin. The core component to warp images is a method to transform an image according to the displacement of a segment.




Figure 5: Example output of the `warpBy1` function on the `bear.png` image using segments: before = `Segment((0,0), (10,0))` and after = `Segment((10, 10), (30, 15))`.

9 Implement `Image warpBy1(const Image &im, Segment segBefore, Segment segAfter)`. This is a resampling function that warps an entire image according to one a pair of segments. **Hint:** Figure 1 of the Beier and Neely paper might be helpful here.

The output should be an image of the same size as the input such that the feature under `segBefore` is now at the location of `segAfter`. Use bilinear reconstruction (with `clamp = true`). Again, use simple examples to test this function, see Figure 5. Once you are done with this, you have completed the hardest part of the assignment.

3.4 Warping according to multiple pairs of segments

In this question, you will extend you warp code to perform transformations according to multiple pairs of segments. For each pixel, transform its 2D coordinates according to each pair of segments and take a weighted average according to the length of **each segment** and the distance of the pixel to the segments. Specifically, the weight is given according to Beier and Neely:

$$\text{weight} = \left(\frac{\text{length}^p}{a + \text{dist}} \right)^b \quad (4)$$

where a, b, p are parameters that control the interpolation. In our test, we have used $b = p = 1$ and $a = 10$ (roughly 5% of the image size).

10.a Implement `float Segment::weight(Vec2f &X, float a, float b, float p)` based on the formula above.

10.b Implement `Image warp(const Image &im, const &vector<Segment> src_segs, const &vector<Segment> dst_segs, float a=10.0, float b=1.0, p=1.0)` which returns a new warped image according

to the list of before and after segments. Pay attention to the order of the segments, where you should loop over output pixels. **Hint:** section 3.3 of the paper and the above function will be helpful here.

Use the provided javascript UI to specify segments. The points must be entered in the same order on the left and right image. You can then copy-paste the C++ code generated below the images to create the corresponding segment objects.

3.5 UI

We provide you with a rudimentary (to say the least) interface to specify segment location from a web browser. It is based on javascript and the raphael library (<http://dmitrybaranovskiy.github.io/raphael/>), and improved by students throughout the years. You specify two input images in `morph-ui.html` (the images must be the same size).

You must click on the segment endpoints in the same order on the left and on the right. Unfortunately, you cannot edit the locations after you have clicked, but you can delete edges by double clicking on them. Once you are done, simply copy the C++ code below each image into your main function to create the corresponding segments.



```
vector<Segment> segsBefore;
segsBefore.push_back(Segment(56, 111, 69, 90));
segsBefore.push_back(Segment(97, 98, 119, 104));
```



```
vector<Segment> segsAfter;
segsAfter.push_back(Segment(61, 103, 78, 76));
segsAfter.push_back(Segment(93, 93, 114, 101));
```

Figure 6: The UI to specify line segments

3.6 Morphing

Given your warping code, we will write a function that generates a morph sequence between two images. Again, make sure you are familiar with morphing from the article.

You are given the source and target images, and a list of segments for each image (the position in the list defines the corresponding pairs of segments, so the

lists should have the same number of elements). You must generate N images morphing from the first input image to the second input image.

For each image, compute its corresponding time fraction t . Then linearly interpolate the position of each segment's endpoints according to t , where $t = 0$ corresponds to the position in the first image and $t = 1$ is the position in the last image. You might want to visualize the results for debugging.

You now need to warp both the first and last image so that their segment locations are moved to the segment location at t , which will align the features of the images. We suggest that you write these two images to disk and verify that the images align and that, as t increases, the images get warped from the configuration of the first image all the way to that of the last one.

Finally, for each t , perform a linear interpolation between the pixel values of the two warped images.

Your function should return a sequence of images. For debugging you can use your main function to write your images to disk using a sequence of names such as `morph_1.png`, `morph_2.png`, ..., see `testMorph` in `a5_main.cpp`.

11 Implement `morph(im0, im1, listSegmentsBefore, listSegmentsAfter, N=1, a=10.0, b=1.0, p=1.0)`. It should return a sequence of N images in addition to the two inputs (i.e., when called with the default value of 1, it only generates one new image for $t = 0.5$). The function should check that `im0` and `im1` have the same size, and throw an exception if not. Note that the interpolation weight should be at the scale of $[0, 1]$, so we need a time constant of $1/(N + 1)$ here.



Figure 7: Morph example. Your morphs will look slightly different depending on what segments you use.

Visualize Results Aside from just looking at the images, you can explore your results at <http://tipix.csail.mit.edu> if you wish. Click on load and load

your morphed images (make sure to respect the required naming). Then explore your results using the pointer, or click "play" from the top-right information panel.

There are several ways to make a video (or a .gif) out of your files (**note** this is not required). You can install ffmpeg <http://ffmpeg.org/> but this involves some number of dependencies. Then use it with, e.g. `ffmpeg -i tes_morph_%02d.png out.gif`.

3.7 Class morph

We'll use the code implemented in this problem set to create a class morph. In order to do that, we need a photo of your face.

12 Submit a selfie (of yourself) as a PNG file named `myface.png` in the **Input** folder of your submission. Feel free to make a silly face if you want. Make sure it's named correctly and in the right folder so that our script can pick up your image.

Make sure your photo follows the following specifications:

- 500 x 600 (500 width, 600 height). There are image rescaling tools online, or you can use what you implemented earlier this pset!
- The background must be a single color of either white or otherwise pale color, and there shouldn't be clutter in your background.
- Your entire face (everything from the neck up) should be visible. Think passport photo.
- The file should be a png. If it is not, there should be plenty of conversion tools online.

In the next pset, we will ask you to run your morphing code from your own face to the face of a peer. Get excited!

4 Extra credit

Here are ideas for extensions you could attempt, for 5% each. At most, on the entire assignment, you can get 10% of extra credit:

- Improve the javascript UI. The amount of extra credit you get is dependent on how much improvement you make. We may also keep your improvements for future classes!
- Extend to movies, where segments are specified at a number of keyframes.
- Morphable face models (see <http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.49.9275>)

5 Submission

Turn in your files to the online submission system and make sure all your files are in the `asst` directory under the root of the zip file. If your code compiles on the submission system, it is organized correctly. The submission system will run code in your main function, but we will not use this code for grading. The submission system should also show you the image your code writes to the `./Output` directory

In the submission system, there will be a form in which you should answer the following questions:

- How long did the assignment take? (in minutes)
- Potential issues with your solution and explanation of partial completion (for partial credit)
- Any extra credit you may have implemented and their function signatures if applicable
- Collaboration acknowledgment (you must write your own code)
- What was most unclear/difficult?
- What was most exciting?